

AI Service Deployment Notebook

This notebook contains steps and code to test, promote, and deploy an Agent as an AI Service.

Note: Notebook code generated using Agent Lab will execute successfully. If code is modified or reordered, there is no guarantee it will successfully execute. For details, see: [Saving your work in Agent Lab as a notebook.](#)

Some familiarity with Python is helpful. This notebook uses Python 3.11.

Contents

This notebook contains the following parts:

1. Setup
2. Initialize all the variables needed by the AI Service
3. Define the AI service function
4. Deploy an AI Service
5. Test the deployed AI Service

1. Set up the environment

Before you can run this notebook, you must perform the following setup tasks:

Connection to WML

This cell defines the credentials required to work with watsonx API for both the execution in the project, as well as the deployment and runtime execution of the function.

Action: Provide the IBM Cloud personal API key. For details, see [documentation](#).

In []:

```
import os
from ibm_watsonx_ai import APIClient, Credentials
import getpass

credentials = Credentials(
    url="https://us-south.ml.cloud.ibm.com",
    api_key=getpass.getpass("Please enter your api key (hit enter): ")
)
```

In []:

```
client = APIClient(credentials)
```

Connecting to a space

A space will be used to host the promoted AI Service.

In []:

```
space_id = "7c318d93-91fe-43e4-9916-5572f3ccd4a8"
client.set.default_space(space_id)
```

Promote asset(s) to space

We will now promote assets we will need to store in the space so that we can access their data from the AI

We will now promote assets we will need to stage in the space so that we can access their data from the AI service.

In []:

```
source_project_id = "76768c52-e794-4361-81ae-c6d303d1a9ad"
```

2. Create the AI service function

We first need to define the AI service function

2.1 Define the function

In []:

```
params = {
    "space_id": space_id,
}

def gen_ai_service(context, params = params, **custom):
    # import dependencies
    from langchain_ibm import ChatWatsonx
    from ibm_watsonx_ai import APIClient
    from ibm_watsonx_ai.foundation_models.utils import Tool, Toolkit
    from langchain_core.messages import AIMessage, HumanMessage
    from langgraph.checkpoint.memory import MemorySaver
    from langgraph.prebuilt import create_react_agent
    import json
    import requests

    model = "meta-llama/llama-3-3-70b-instruct"

    service_url = "https://us-south.ml.cloud.ibm.com"
    # Get credentials token
    credentials = {
        "url": service_url,
        "token": context.generate_token()
    }

    # Setup client
    client = APIClient(credentials)
    space_id = params.get("space_id")
    client.set.default_space(space_id)

    def create_chat_model(watsonx_client):
        parameters = {
            "frequency_penalty": 0,
            "max_tokens": 200,
            "presence_penalty": 0,
            "temperature": 0,
            "top_p": 1
        }

        chat_model = ChatWatsonx(
            model_id=model,
            url=service_url,
            space_id=space_id,
            params=parameters,
            watsonx_client=watsonx_client,
        )
        return chat_model

    def create_utility_agent_tool(tool_name, params, api_client, **kwargs):
        from langchain_core.tools import StructuredTool
        utility_agent_tool = Toolkit(
            api_client=api_client
```

```

).get_tool(tool_name)

tool_description = utility_agent_tool.get("description")

if (kwargs.get("tool_description")):
    tool_description = kwargs.get("tool_description")
elif (utility_agent_tool.get("agent_description")):
    tool_description = utility_agent_tool.get("agent_description")

tool_schema = utility_agent_tool.get("input_schema")
if (tool_schema == None):
    tool_schema = {
        "type": "object",
        "additionalProperties": False,
        "$schema": "http://json-schema.org/draft-07/schema#",
        "properties": {
            "input": {
                "description": "input for the tool",
                "type": "string"
            }
        }
    }

def run_tool(**tool_input):
    query = tool_input
    if (utility_agent_tool.get("input_schema") == None):
        query = tool_input.get("input")

    results = utility_agent_tool.run(
        input=query,
        config=params
    )

    return results.get("output")

return StructuredTool(
    name=tool_name,
    description = tool_description,
    func=run_tool,
    args_schema=tool_schema
)

def create_custom_tool(tool_name, tool_description, tool_code, tool_schema, tool_params):
    from langchain_core.tools import StructuredTool
    import ast

    def call_tool(**kwargs):
        tree = ast.parse(tool_code, mode="exec")
        custom_tool_functions = [ x for x in tree.body if isinstance(x, ast.FunctionDef) ]

        function_name = custom_tool_functions[0].name
        compiled_code = compile(tree, 'custom_tool', 'exec')
        namespace = tool_params if tool_params else {}
        exec(compiled_code, namespace)
        return namespace[function_name](**kwargs)

    tool = StructuredTool(
        name=tool_name,
        description = tool_description,
        func=call_tool,
        args_schema=tool_schema
    )
    return tool

def create_custom_tools():
    custom_tools = []

def create_tools(inner_client, context):
    tools = []

```

```

config = None
tools.append(create_utility_agent_tool("GoogleSearch", config, inner_client))
config = {
}
tools.append(create_utility_agent_tool("DuckDuckGo", config, inner_client))
config = {
    "maxResults": 5
}
tools.append(create_utility_agent_tool("Wikipedia", config, inner_client))
config = {
}
tools.append(create_utility_agent_tool("Weather", config, inner_client))
config = {
}
tools.append(create_utility_agent_tool("WebCrawler", config, inner_client))
return tools

```

```

def create_agent(model, tools, messages):
    memory = MemorySaver()
    instructions = """# Notes

```

- Use markdown syntax for formatting code snippets, links, JSON, tables, images, files.
- Any HTML tags must be wrapped in block quotes, for example ``<html>``.
- When returning code blocks, specify language.
- Sometimes, things don't go as planned. Tools may not provide useful information on the first few tries. You should always try a few different approaches before declaring the problem unsolvable.
- When the tool doesn't give you what you were asking for, you must either use another tool or a different tool input.
- When using search engines, you try different formulations of the query, possibly even in a different language.
- You cannot do complex calculations, computations, or data manipulations without using tools.
- If you need to call a tool to compute something, always call it instead of saying you will call it.

If a tool returns an IMAGE in the result, you must include it in your answer as Markdown.

Example:

```

Tool result: IMAGE({commonApiUrl}/wx/v1-beta/utility_agent_tools/cache/images/plt-04e3c91ae04b47f8934a4e6b7d1fdc2c.png)
Markdown to return to user: ![Generated image]({commonApiUrl}/wx/v1-beta/utility_agent_tools/cache/images/plt-04e3c91ae04b47f8934a4e6b7d1fdc2c.png)

```

You are a multilingual AI health assistant. A user will describe their symptoms in natural language.

Your job is to:

1. List possible conditions (based on WHO and trusted medical sources).
2. Rate the urgency: Low, Medium, High, Emergency.
3. Give home remedies and preventive care.
4. Tell if a doctor should be consulted.
5. Provide this info in user's preferred language.

Never give direct medical diagnosis. Avoid treatment instructions that require prescriptions. Encourage users to seek licensed healthcare professionals.

Respond in bullet format. Be polite, professional, and helpful.

```

"""
    for message in messages:
        if message["role"] == "system":
            instructions += message["content"]
    graph = create_react_agent(model, tools=tools, checkpoint=memory, state_modifier=instructions)
    return graph

```

```

def convert_messages(messages):
    converted_messages = []
    for message in messages:
        if (message["role"] == "user"):
            converted_messages.append(HumanMessage(content=message["content"]))
        elif (message["role"] == "assistant"):

```

```

        converted_messages.append(AIMessage(content=message["content"]))
    return converted_messages

def generate(context):
    payload = context.get_json()
    messages = payload.get("messages")
    inner_credentials = {
        "url": service_url,
        "token": context.get_token()
    }

    inner_client = APIClient(inner_credentials)
    model = create_chat_model(inner_client)
    tools = create_tools(inner_client, context)
    agent = create_agent(model, tools, messages)

    generated_response = agent.invoke(
        { "messages": convert_messages(messages) },
        { "configurable": { "thread_id": "42" } }
    )

    last_message = generated_response["messages"][-1]
    generated_response = last_message.content

    execute_response = {
        "headers": {
            "Content-Type": "application/json"
        },
        "body": {
            "choices": [{
                "index": 0,
                "message": {
                    "role": "assistant",
                    "content": generated_response
                }
            }]
        }
    }

    return execute_response

def generate_stream(context):
    print("Generate stream", flush=True)
    payload = context.get_json()
    headers = context.get_headers()
    is_assistant = headers.get("X-Ai-Interface") == "assistant"
    messages = payload.get("messages")
    inner_credentials = {
        "url": service_url,
        "token": context.get_token()
    }

    inner_client = APIClient(inner_credentials)
    model = create_chat_model(inner_client)
    tools = create_tools(inner_client, context)
    agent = create_agent(model, tools, messages)

    response_stream = agent.stream(
        { "messages": messages },
        { "configurable": { "thread_id": "42" } },
        stream_mode=["updates", "messages"]
    )

    for chunk in response_stream:
        chunk_type = chunk[0]
        finish_reason = ""
        usage = None
        if (chunk_type == "messages"):
            message_object = chunk[1][0]
            if (message_object.type == "AIMessageChunk" and message_object.content != ""):
                message = {
                    "role": "assistant",

```

```

        "content": message_object.content
    }
    else:
        continue
elif (chunk_type == "updates"):
    update = chunk[1]
    if ("agent" in update):
        agent = update["agent"]
        agent_result = agent["messages"][0]
        if (agent_result.additional_kwargs):
            kwargs = agent["messages"][0].additional_kwargs
            tool_call = kwargs["tool_calls"][0]
            if (is_assistant):
                message = {
                    "role": "assistant",
                    "step_details": {
                        "type": "tool_calls",
                        "tool_calls": [
                            {
                                "id": tool_call["id"],
                                "name": tool_call["function"]["name"],
                                "args": tool_call["function"]["arguments"]
                            }
                        ]
                    }
                }
            else:
                message = {
                    "role": "assistant",
                    "tool_calls": [
                        {
                            "id": tool_call["id"],
                            "type": "function",
                            "function": {
                                "name": tool_call["function"]["name"],
                                "arguments": tool_call["function"]["argument
s"]
                            }
                        }
                    ]
                }
        elif (agent_result.response_metadata):
            # Final update
            message = {
                "role": "assistant",
                "content": agent_result.content
            }
            finish_reason = agent_result.response_metadata["finish_reason"]
            if (finish_reason):
                message["content"] = ""

            usage = {
                "completion_tokens": agent_result.usage_metadata["output_tok
ens"],
                "prompt_tokens": agent_result.usage_metadata["input_tokens"]
            ,
                "total_tokens": agent_result.usage_metadata["total_tokens"]
            }
        elif ("tools" in update):
            tools = update["tools"]
            tool_result = tools["messages"][0]
            if (is_assistant):
                message = {
                    "role": "assistant",
                    "step_details": {
                        "type": "tool_response",
                        "id": tool_result.id,
                        "tool_call_id": tool_result.tool_call_id,
                        "name": tool_result.name,
                        "content": tool_result.content
                    }
                }

```

```

        else:
            message = {
                "role": "tool",
                "id": tool_result.id,
                "tool_call_id": tool_result.tool_call_id,
                "name": tool_result.name,
                "content": tool_result.content
            }
        else:
            continue

    chunk_response = {
        "choices": [{
            "index": 0,
            "delta": message
        }]
    }
    if (finish_reason):
        chunk_response["choices"][0]["finish_reason"] = finish_reason
    if (usage):
        chunk_response["usage"] = usage
    yield chunk_response

return generate, generate_stream

```

2.2 Test locally

In []:

```

# Initialize AI Service function locally
from ibm_watsonx_ai.deployments import RuntimeContext

context = RuntimeContext(api_client=client)

streaming = False
findex = 1 if streaming else 0
local_function = gen_ai_service(context, space_id=space_id)[findex]
messages = []

```

In []:

```

local_question = "Change this question to test your function"

messages.append({ "role" : "user", "content": local_question })

context = RuntimeContext(api_client=client, request_payload_json={"messages": messages})

response = local_function(context)

result = ''

if (streaming):
    for chunk in response:
        print(chunk, end="\n\n", flush=True)
else:
    print(response)

```

3. Store and deploy the AI Service

Before you can deploy the AI Service, you must store the AI service in your watsonx.ai repository.

In []:

```

# Look up software specification for the AI service
software_spec_id_in_project = "45f12dfe-aa78-5b8d-9f38-0ee223c47309"
software_spec_id = ""

try:

```

```

software_spec_id = client.software_specifications.get_id_by_name("runtime-24.1-py3.11")
except:
    software_spec_id = client.spaces.promote(software_spec_id_in_project, source_project_id, space_id)

```

In []:

```

# Define the request and response schemas for the AI service
request_schema = {
    "application/json": {
        "$schema": "http://json-schema.org/draft-07/schema#",
        "type": "object",
        "properties": {
            "messages": {
                "title": "The messages for this chat session.",
                "type": "array",
                "items": {
                    "type": "object",
                    "properties": {
                        "role": {
                            "title": "The role of the message author.",
                            "type": "string",
                            "enum": ["user", "assistant"]
                        },
                        "content": {
                            "title": "The contents of the message.",
                            "type": "string"
                        }
                    }
                },
                "required": ["role", "content"]
            }
        },
        "required": ["messages"]
    }
}

response_schema = {
    "application/json": {
        "oneOf": [{"$schema": "http://json-schema.org/draft-07/schema#", "type": "object", "description": "AI Service response for /ai_service_stream", "properties": {"choices": {"description": "A list of chat completion choices.", "type": "array", "items": {"type": "object", "properties": {"index": {"type": "integer", "title": "The index of this result."}, "delta": {"description": "A message result.", "type": "object", "properties": {"content": {"description": "The contents of the message.", "type": "string"}, "role": {"description": "The role of the author of this message.", "type": "string"}}, "required": ["role"]}}}}, {"$schema": "http://json-schema.org/draft-07/schema#", "type": "object", "description": "AI Service response for /ai_service", "properties": {"choices": {"description": "A list of chat completion choices", "type": "array", "items": {"type": "object", "properties": {"index": {"type": "integer", "description": "The index of this result."}, "message": {"description": "A message result.", "type": "object", "properties": {"role": {"description": "The role of the author of this message.", "type": "string"}, "content": {"title": "Message content.", "type": "string"}}, "required": ["role"]}}}}, {"required": ["choices"]}]
    }
}

```

In []:

```

# Store the AI service in the repository
ai_service_metadata = {
    client.repository.AIServiceMetaNames.NAME: "symptoms",
    client.repository.AIServiceMetaNames.DESCRPTION: "",
    client.repository.AIServiceMetaNames.SOFTWARE_SPEC_ID: software_spec_id,
    client.repository.AIServiceMetaNames.CUSTOM: {},
    client.repository.AIServiceMetaNames.REQUEST_DOCUMENTATION: request_schema,
    client.repository.AIServiceMetaNames.RESPONSE_DOCUMENTATION: response_schema,
    client.repository.AIServiceMetaNames.TAGS: ["wx-agent"]
}

ai_service_details = client.repository.store_ai_service(meta_props=ai_service_metadata, a

```



```
i_service=gen_ai_service)
```

```
In [ ]:
```

```
# Get the AI Service ID
```

```
ai_service_id = client.repository.get_ai_service_id(ai_service_details)
```

```
In [ ]:
```

```
# Deploy the stored AI Service
```

```
deployment_custom = {  
    "avatar_icon": "Chemistry",  
    "avatar_color": "supportCautionMajor",  
    "placeholder_image": "placeholder2.png"  
}  
deployment_metadata = {  
    client.deployments.ConfigurationMetaNames.NAME: "symptoms",  
    client.deployments.ConfigurationMetaNames.ONLINE: {},  
    client.deployments.ConfigurationMetaNames.CUSTOM: deployment_custom,  
    client.deployments.ConfigurationMetaNames.DESCRPTION: "predicts the disease by symptoms",  
    client.repository.AIServiceMetaNames.TAGS: ["wx-agent"]  
}
```

```
function_deployment_details = client.deployments.create(ai_service_id, meta_props=deployment_metadata, space_id=space_id)
```

4. Test AI Service

```
In [ ]:
```

```
# Get the ID of the AI Service deployment just created
```

```
deployment_id = client.deployments.get_id(function_deployment_details)  
print(deployment_id)
```

```
In [ ]:
```

```
messages = []  
remote_question = "Change this question to test your function"  
messages.append({ "role" : "user", "content": remote_question })  
payload = { "messages": messages }
```

```
In [ ]:
```

```
result = client.deployments.run_ai_service(deployment_id, payload)  
if "error" in result:  
    print(result["error"])  
else:  
    print(result)
```

Next steps

You successfully deployed and tested the AI Service! You can now view your deployment and test it as a REST API endpoint.

Copyrights

Licensed Materials - Copyright © 2024 IBM. This notebook and its source code are released under the terms of the ILAN License. Use, duplication disclosure restricted by GSA ADP Schedule Contract with IBM Corp.

Note: The auto-generated notebooks are subject to the International License Agreement for Non-Warranted Programs (or equivalent) and License Information document for watsonx.ai Auto-generated Notebook (License Terms), such agreements located in the link below. Specifically, the Source Components and Sample Materials clause included in the License Information document for watsonx.ai Studio Auto-generated Notebook applies to

the auto-generated notebooks.

By downloading, copying, accessing, or otherwise using the materials, you agree to the [License Terms](#)